




Two-model comparison on challenge v0.6 post-regression (n = 60)

DeepSeek v4 Flash GPT-5.4-mini*



* GPT model identity is asserted by a third-party relay. Token counts are provider-reported and may not be directly comparable. Single temperature-0 runs; descriptive results without confidence intervals.